

Google Cloud Professional Machine Learning Engineer Certification: Course Outline

No of Days – 5 (40 Hours)

Section 1: Framing ML problems

1.1 Translating business challenges into ML use cases. Considerations include:

- Choosing the best solution (ML vs. non-ML, custom vs. pre-packaged [e.g., AutoML, Vision API]) based on the business requirements
- Defining how the model output should be used to solve the business problem
- Deciding how incorrect results should be handled
- Identifying data sources (available vs. ideal)

1.2 Defining ML problems. Considerations include:

- Problem type (e.g., classification, regression, clustering)
- Outcome of model predictions
- Input (features) and predicted output format

1.3 Defining business success criteria. Considerations include:

- Alignment of ML success metrics to the business problem
- Key results
- Determining when a model is deemed unsuccessful

1.4 Identifying risks to feasibility of ML solutions. Considerations include:

- Assessing and communicating business impact
- Assessing ML solution readiness
- Assessing data readiness and potential limitations
- Aligning with Google's Responsible AI practices (e.g., different biases)

Section 2: Architecting ML solutions

2.1 Designing reliable, scalable, and highly available ML solutions. Considerations include:

- Choosing appropriate ML services for the use case (e.g., Cloud Build, Kubeflow)
- Component types (e.g., data collection, data management)
- Exploration/analysis
- Feature engineering
- Logging/management
- Automation
- Orchestration

- Monitoring
- Serving

2.2 Choosing appropriate Google Cloud hardware components. Considerations include:

- Evaluation of compute and accelerator options (e.g., CPU, GPU, TPU, edge devices)

2.3 Designing architecture that complies with security concerns across sectors/industries.

Considerations include:

- Building secure ML systems (e.g., protecting against unintentional exploitation of data/model, hacking)
- Privacy implications of data usage and/or collection (e.g., handling sensitive data such as Personally Identifiable Information [PII] and Protected Health Information [PHI])

Section 3: Designing data preparation and processing systems

3.1 Exploring data (EDA). Considerations include:

- Visualization
- Statistical fundamentals at scale
- Evaluation of data quality and feasibility
- Establishing data constraints (e.g., TFDV)

3.2 Building data pipelines. Considerations include:

- Organizing and optimizing training datasets
- Data validation
- Handling missing data
- Handling outliers
- Data leakage

3.3 Creating input features (feature engineering). Considerations include:

- Ensuring consistent data pre-processing between training and serving
- Encoding structured data types
- Feature selection
- Class imbalance
- Feature crosses
- Transformations (TensorFlow Transform)

Section 4: Developing ML models

4.1 Building models. Considerations include:

- Choice of framework and model
- Modeling techniques given interpretability requirements
- Transfer learning
- Data augmentation
- Semi-supervised learning
- Model generalization and strategies to handle overfitting and underfitting

4.2 Training models. Considerations include:

- Ingestion of various file types into training (e.g., CSV, JSON, IMG, parquet or databases, Hadoop/Spark)
- Training a model as a job in different environments
- Hyperparameter tuning
- Tracking metrics during training
- Retraining/redeployment evaluation

4.3 Testing models. Considerations include:

- Unit tests for model training and serving
- Model performance against baselines, simpler models, and across the time dimension
- Model explainability on Vertex AI

4.4 Scaling model training and serving. Considerations include:

- Distributed training
- Scaling prediction service (e.g., Vertex AI Prediction, containerized serving)

Section 5: Automating and orchestrating ML pipelines

5.1 Designing and implementing training pipelines. Considerations include:

- Identification of components, parameters, triggers, and compute needs (e.g., Cloud Build, Cloud Run)
- Orchestration framework (e.g., Kubeflow Pipelines/Vertex AI Pipelines, Cloud Composer/Apache Airflow)
- Hybrid or multicloud strategies
- System design with TFX components/Kubeflow DSL

5.2 Implementing serving pipelines. Considerations include:

- Serving (online, batch, caching)
- Google Cloud serving options
- Testing for target performance
- Configuring trigger and pipeline schedules

5.3 Tracking and auditing metadata. Considerations include:

- Organizing and tracking experiments and pipeline runs
- Hooking into model and dataset versioning
- Model/dataset lineage

Section 6: Monitoring, optimizing, and maintaining ML solutions

6.1 Monitoring and troubleshooting ML solutions. Considerations include:

- Performance and business quality of ML model predictions
- Logging strategies
- Establishing continuous evaluation metrics (e.g., evaluation of drift or bias)
- Understanding Google Cloud permissions model
- Identification of appropriate retraining policy
- Common training and serving errors (TensorFlow)
- ML model failure and resulting biases

6.2 Tuning performance of ML solutions for training and serving in production.

Considerations include:

- Optimization and simplification of input pipeline for training
- Simplification techniques

Certification Details:

- This certification is of 2 hours duration for completing the 60 questions
- Questions within this exams are of unique scenario
- It is available in two languages i.e. English and Japanese
- All the questions are of Multiple Choice and multiple select taken remotely or in person at a test center
- The exam can be delivered in two modes: Take the online-proctored exam from a remote location, review the online testing [requirements](#)
- Take the onsite-proctored exam at a testing center, [Locate a test center near you](#).